

- 1.2.a) Refer to the documentation, what is the functionality of the `tol` parameter in the Perceptron class? (2 marks)

`tol` represents the improvement in the model's error. The model will return a solution when the improvement in error between two iterations is smaller than `tol`, i.e. the improvement is considered small and a good solution has been found.

- 1.2.b) If we set `max_iter=5000` and `tol=1e-3` (the rest as default), does this guarantee that the algorithm will pass over the training data 5000 times? If not, ensure that the algorithm will pass over the training data 5000 times? (2 marks)

No it does not. To ensure that the algorithm will pass over the training data 5000 times, set `tol` to 0.

- 1.2.c) How can we set the weights of the model to a certain value? (2 marks)

First fit the model. Then modify the `coef_` and `intercept_` attribute manually. The `coef_` attribute sets the weight for each feature, and the `intercept_` sets the bias term.

- 1.2.d) How close is the performance (through confusion matrix) of your NumPy implementation in comparison to the existing modules in the scikit-learn library? (2 marks)

My implementation's confusion matrix:

```
[[8. 0.]
 [3. 9.]]
```

Scikit-learn library's confusion matrix:

```
[[8. 0.]
 [2. 10.]]
```

My implementation misclassified one more point in class +1 than the library's module.

- 2.1.a) When we input a singular matrix, the function `linalg.inv` often returns an error message. In your `fit_LinRegr(X_train, y_train)` implementation, is your input to the function `linalg.inv` a singular matrix? Explain why. (2 marks)

The input is indeed a singular matrix because its determinant is 0, making it non-invertible.

- 2.1.b) As you are using `linalg.inv` for matrix inversion, report the output message when running the function `subtestFn()`. We note that inputting a singular matrix to `linalg.inv` sometimes does not yield an error due to numerical issue. (1 marks)

It outputs ERROR.

- 2.1.c) Replace the function `linalg.inv` with `linalg.pinv`, you should get the model's weight and the "NO ERROR" message after running the function `subtestFn()`. Explain the difference between `linalg.inv` and `linalg.pinv`, and report the model's weight. (2 marks)

`linalg.pinv` performs pseudo inverse, which can handle both singular and non-singular matrices. `linalg.inv` is used for non-singular matrices.

The weight is: [1.04360964e-14, 2.00000000e-01, 4.00000000e-01]